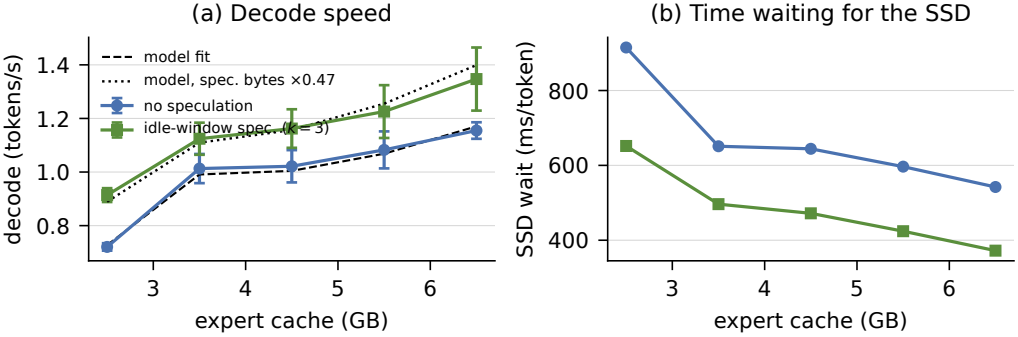
(a) Decode speed
decode (tokens/s)
0.8
1.0
1.2
1.4
model fit
model, spec. bytes ×0.47
no speculation
idle-window spec. (K = 3)
3
4
5
6
expert cache (GB)
(b) Time waiting for the SSD
SSD wait (ms/token)
400
600
800
3
4
5
6
expert cache (GB)