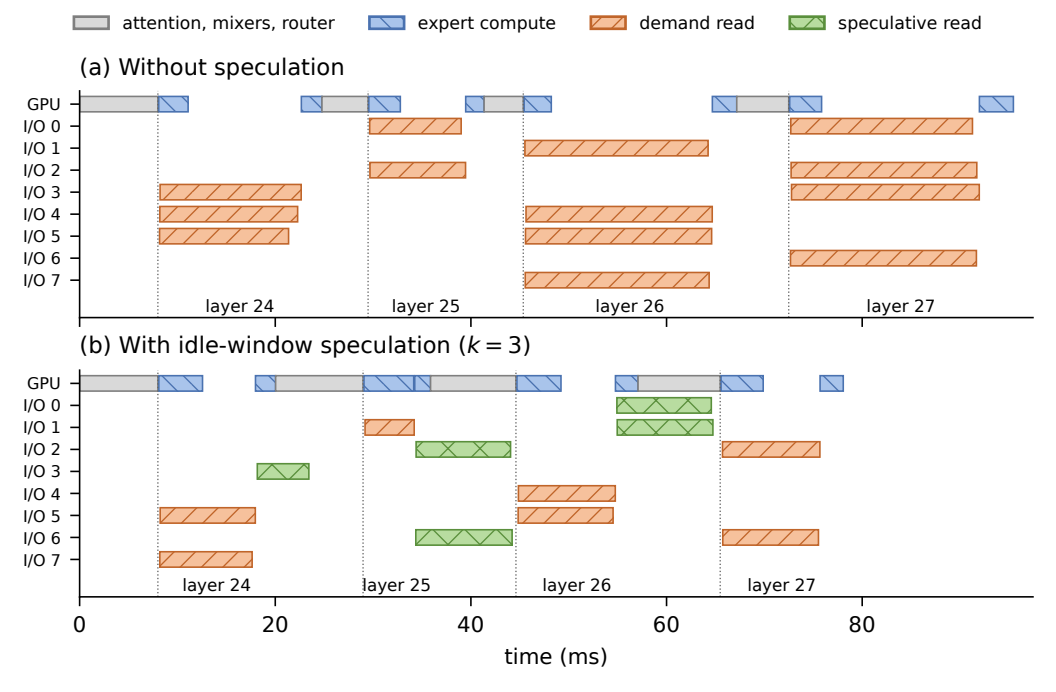
attention, mixers, router
expert compute
demand read
speculative read
(a) Without speculation
GPU
I/O 0
I/O 1
I/O 2
I/O 3
I/O 4
I/O 5
I/O 6
I/O 7
layer 24
layer 25
layer 26
layer 27
(b) With idle-window speculation ($k = 3$)
GPU
I/O 0
I/O 1
I/O 2
I/O 3
I/O 4
I/O 5
I/O 6
I/O 7
layer 24
layer 25
layer 26
layer 27
0
20
40
60
80
time (ms)